

validate_model

April 16, 2025

```
[18]: print("hi")
```

hi

```
[19]: from ultralytics import YOLO
import os
%matplotlib inline
import matplotlib.pyplot as plt
from sklearn.metrics import confusion_matrix, classification_report
import seaborn as sns
import numpy as np
```

```
[20]: # Load the YOLO model
base_dir = '/home/ian/intellicook/ingredient-recognition/model'
weights_path = os.path.join(base_dir, 'train/runs/detect/train/weights/best.pt')
model = YOLO(weights_path)
```

```
[21]: # Validate the model
data_path = os.path.join(base_dir, 'data/Food Ingredient Recognition.v4i.
↳ yolov11/data.yaml')
results = model.val(data=data_path, split='val')
```

Ultralytics 8.3.55 Python-3.12.2 torch-2.6.0+cu124 CUDA:0 (NVIDIA GeForce RTX 4090, 24195MiB)

YOLO11m summary (fused): 303 layers, 20,077,063 parameters, 0 gradients, 67.9 GFLOPs

val: Scanning /home/ian/intellicook/ingredient-recognition/model/data/Food Ingredient

Recognition.v4i.yolov11/valid/labels.cache... 780 images, 67 backgrounds, 0 corrupt: 100%|

	Class	Images	Instances	Box(P	R	mAP50
mAP50-95): 100%		49/49	[00:04<00:00, 11.94it/s]			
0.594	all	780	2517	0.8	0.815	0.858
0.81	apple	10	20	0.947	0.9	0.919
	asparagus	12	47	0.532	0.298	0.371

0.117	avocado	13	49	0.934	0.871	0.963
0.709	banana	15	25	0.768	0.92	0.878
0.633	bell pepper	13	59	0.851	0.972	0.966
0.544	bitter gourd	10	33	0.581	0.606	0.625
0.258	bok choy	17	49	0.752	0.682	0.796
0.337	broccoli	7	27	0.833	0.889	0.886
0.498	cabbage	10	19	0.864	0.947	0.969
0.674	carrot	14	39	0.705	0.538	0.672
0.267	cashew	14	94	0.836	0.921	0.92
0.434	cauliflower	14	17	0.932	0.882	0.983
0.576	chayote	10	51	0.87	0.961	0.953
0.545	chicken breast	20	38	0.909	0.895	0.97
0.633	chicken thigh	18	31	0.979	0.935	0.978
0.741	chicken wing	13	34	0.859	0.735	0.866
0.531	chilli	13	55	0.53	0.709	0.63
0.41	coconut	19	49	0.709	0.918	0.895
0.742	coconuts	3	10	0.93	0.9	0.978
0.642	corn	18	79	0.736	0.835	0.822
0.501	crab	11	18	0.511	0.444	0.515
0.316	cucumber	15	71	0.815	0.808	0.899
0.562	egg_	13	61	0.977	0.934	0.97
0.861	eggplant	16	60	0.785	0.75	0.777
0.573	garlic	13	53	0.947	0.906	0.962
0.731	ginger	6	14	0.675	0.5	0.568

0.421						
	grapes	11	20	0.764	0.75	0.84
0.584						
	lemon	26	65	0.815	0.678	0.848
0.657						
	lettuce	10	13	0.723	0.604	0.628
0.424						
	lobster tails	11	28	0.832	0.786	0.887
0.539						
	mango	14	47	0.856	0.809	0.937
0.818						
	melon	20	41	0.864	0.854	0.888
0.665						
	onion	8	53	0.596	1	0.933
0.726						
	orange	10	36	0.871	0.917	0.98
0.768						
	oysters	12	87	0.891	0.862	0.878
0.658						
	pawpaw	11	33	0.691	0.909	0.841
0.631						
	peanuts	20	222	0.726	0.824	0.863
0.584						
	peas	7	12	0.704	0.667	0.73
0.348						
	pineapple	17	27	0.924	0.815	0.893
0.657						
	pork belly	13	29	0.692	0.759	0.841
0.62						
	potato	16	71	0.922	0.986	0.988
0.868						
	pumpkin	11	19	0.849	0.947	0.965
0.87						
	radishes	13	19	0.713	0.785	0.793
0.436						
	red rice	14	23	0.79	0.826	0.94
0.741						
	salmon	17	27	0.884	0.848	0.954
0.695						
	sea scallops	6	43	0.761	0.907	0.86
0.651						
	shrimp	10	47	0.717	0.66	0.674
0.409						
	spinach	20	156	0.71	0.763	0.806
0.504						
	strawberry	12	59	0.79	0.814	0.904
0.591						
	sweet potato	19	77	0.865	0.922	0.957

0.754	tempeh	11	34	0.788	0.877	0.854
0.503	tofu	15	58	0.863	0.862	0.942
0.704	tomato	15	43	0.75	0.93	0.877
0.773	tuna	10	16	0.936	0.812	0.953
0.702	white rice	8	10	0.948	1	0.995

0.695

Speed: 0.3ms preprocess, 3.5ms inference, 0.0ms loss, 0.3ms postprocess per image

Results saved to runs/detect/val2

```
[22]: # Extract metrics
metrics = results.results_dict
print(metrics.keys())

for key in metrics:
    print(f"{key}: {metrics[key]}")
```

```
dict_keys(['metrics/precision(B)', 'metrics/recall(B)', 'metrics/mAP50(B)',
'metrics/mAP50-95(B)', 'fitness'])
metrics/precision(B): 0.8000694712113787
metrics/recall(B): 0.8150705497445575
metrics/mAP50(B): 0.8577925524501078
metrics/mAP50-95(B): 0.5935191361677661
fitness: 0.6199464777960002
```

```
[23]: image_paths = [
    os.path.join(base_dir, 'train/runs/detect/val/confusion_matrix.png'),
    os.path.join(base_dir, 'train/runs/detect/val/F1_curve.png'),
    os.path.join(base_dir, 'train/runs/detect/val/PR_curve.png'),
    os.path.join(base_dir, 'train/runs/detect/val/P_curve.png'),
    os.path.join(base_dir, 'train/runs/detect/val/R_curve.png'),
    os.path.join(base_dir, 'train/runs/detect/val/confusion_matrix_normalized.
    png'),
    os.path.join(base_dir, 'train/runs/detect/val/val_batch0_labels.jpg'),
    os.path.join(base_dir, 'train/runs/detect/val/val_batch0_pred.jpg')
]

# Display all images
fig, axes = plt.subplots(nrows=4, ncols=2, figsize=(15, 20))
axes = axes.flatten()

for ax, img_path in zip(axes, image_paths):
```

```
img = plt.imread(img_path)
ax.imshow(img)
ax.axis('off') # Hide the axis
ax.set_title(os.path.basename(img_path))

plt.tight_layout()
plt.show()
```

